

KAMLENDU KUMAR

Data Engineer | AWS · Spark Structured Streaming · Apache Iceberg · Databricks

Pune, India • Open to remote (IST, overlap to EU / US-East) • +91 62033 20570 • kamlendukumar4480@gmail.com

SUMMARY

Data engineer with 3+ years at Barclays building AWS-native streaming and lakehouse infrastructure for the bank's centralized customer master-data platform. Owns near-real-time Spark Structured Streaming pipelines, an Apache Iceberg lakehouse, and the on-premise-Hadoop-to-AWS migration that moved it all — under KYC/AML, tokenization, and disaster-recovery constraints. AWS-certified across four tracks.

TECHNICAL SKILLS

Streaming & Processing Apache Spark, PySpark, Spark Structured Streaming, Hadoop, near-real-time ingestion

Lakehouse & Warehousing Apache Iceberg, Delta Lake, Databricks, Snowflake, Amazon Redshift, table maintenance & compaction

AWS S3, Glue, EMR, Lambda, Redshift, Service Catalog, IAM, disaster recovery & resiliency

Pipelines & Migration ETL/ELT design, Ab Initio, cloud migration & data hydration, entity resolution, deduplication, data quality

Governance & Interfaces Data governance, PII/CII/CALT tokenization, KYC/AML regulatory reporting, RESTful APIs

Languages Python, SQL, C++, C#

EXPERIENCE

Cloud Data Engineer — Barclays Jul 2023 – Present · Pune, India

Open Party Platform (OPP) — the bank's centralized party/customer master-data platform and single source of truth across business lines.

- **Build and operate near-real-time streaming ingestion** on Spark Structured Streaming, delivering low-latency party updates and cross-platform synchronization to downstream systems bank-wide.
- **Led the migration and hydration of the platform from on-premise Hadoop to AWS**, including disaster-recovery and resiliency server setup, sustaining business continuity and high availability through cutover.
- **Own the Apache Iceberg lakehouse layer** — automated compaction and deletion, table maintenance, and AWS-native Databricks infrastructure design that reduced compute spend.
- **Ship ingestion and preprocessing pipelines** behind real-time Party Search and Party Detail Retrieval REST APIs consumed by KYC, AML screening, onboarding, and regulatory-reporting systems.
- **Built entity-matching and deduplication pipelines** that identify and merge duplicate party records, measurably improving data quality and integrity across business lines.
- **Implemented PII/CII/CALT tokenization with Ab Initio** to meet regulatory obligations, and owned data-governance and risk stories end to end with non-technical stakeholders.
- **Authored 30+ AWS Service Catalog products** (including governed Redshift clusters), enabling compliant self-service infrastructure provisioning for multiple engineering teams.

SELECTED PROJECTS

Attestic — Migration Assurance Platform Node.js · Ed25519 · multi-warehouse connectors

Reconciliation platform producing deterministic, cryptographically signed evidence bundles for cloud data migrations, with connectors across Oracle, SQL Server, Db2, Snowflake, Databricks, BigQuery, and Iceberg/Delta.

Career Autopilot — AI Job-Search Platform React · Python · LLM APIs

Parses resumes against live postings, scores fit via an LLM, and generates tailored application materials and platform-specific Boolean search strings; signed proof-of-work credentials on a microservices backend.

Bharat Radar — Options Analytics Python · JavaScript · LLM APIs

Black-Scholes pricing with full Greeks for NSE F&O instruments, plus a model-evaluation framework grading trade setups with confidence tiers and web-search augmentation.

EDUCATION

B.Tech, Computer Science & Engineering (Cloud Computing & Virtualization) · UPES Dehradun, India 2019 – 2023

CERTIFICATIONS

AWS Solutions Architect – Associate · Generative AI Developer – Professional · AI Practitioner · Cloud Practitioner

Other Snowflake Pro Associate · Oracle Certified Expert (SOA) · Databricks Certified Data Engineer Associate (in progress)